


**Exercise: a)** Load data and visualise it using gscatter:

```
load('classification.mat');
gscatter(X(:,1), X(:,2), y, 'rb', 'o+');
xlabel('x^{(1)}'), ylabel('x^{(2)}');
```



**Exercise: b)** Use holdout function to split data set into training set and test set with some `trainRatio` while keeping 0/1s ratio of original data set in both. `trainRatio = 0.6` should ensure sufficient accuracy later on the test set.

```
function [X_train, X_test, y_train, y_test] = holdout(X, y, trainRatio)
N = size(X, 1);
y_0 = find(y == 0); %Array of sample entries x_i with y = 0
T_0 = datasample(y_0, floor(trainRatio*size(y_0,1)), 'Replace', false);
y_1 = find(y); %Array of sample entries x_i with y = 1
T_1 = datasample(y_1, floor(trainRatio*size(y_0,1)), 'Replace', false);

ind_train = [T_0; T_1]; ind_test = setdiff(1:N, ind_train); %indices
X_train = X(ind_train, :); y_train = y(ind_train);
X_test = X(ind_test, :); y_test = y(ind_test);
end
```

**Exercise: c)** Consider a Logistic Classification model for extended inputs  $\tilde{x}_n = (1, x_n)^T$  for  $n = 1, \dots, N$  with parameter  $\beta$ :

$$P(y_n = 1 | \tilde{x}_n) = \frac{1}{1 + \exp(-\beta^T \tilde{x}_n)} := \sigma(\beta^T \tilde{x}_n)$$

$$\text{Loglikelihood: } \mathcal{L}(\beta) = \sum_{n=1}^N y_n \log(\sigma(\beta^T \tilde{x}_n)) + (1 - y_n) \log(1 - \sigma(\beta^T \tilde{x}_n))$$

Gradient of loglikelihood using derivative identity of  $\sigma$ :

$$\frac{\partial \mathcal{L}(\beta)}{\partial \beta} = \sum_{n=1}^N y_n \frac{(1 - \sigma(\beta^T \tilde{x}_n)) \sigma(\beta^T \tilde{x}_n)}{\sigma(\beta^T \tilde{x}_n)} \tilde{x}_n - (1 - y_n) \frac{(1 - \sigma(\beta^T \tilde{x}_n)) \sigma(\beta^T \tilde{x}_n)}{1 - \sigma(\beta^T \tilde{x}_n)} \tilde{x}_n = \sum_{n=1}^N [y_n - \sigma(\beta^T \tilde{x}_n)] \tilde{x}_n$$

We compute the loglikelihood for  $\beta$  and its gradient given any extended data set `X_ext` and labels `y` with:


```
function [loglike, grad] = log_likelihood_grad(X_ext, y, beta)
N = size(X_ext, 1);
%My function for Logistic classification model P(y=1|x_ext) = sigma(beta'x_ext)
sigma = probs_logreg(X_ext, beta);
%Initialise output variables
loglike = 0; grad=0;
%Term by term sum
for n=1:N
    sum_term = y(n)*log(sigma(n)) + (1 - y(n))*log(1 - sigma(n));
    loglike = loglike + sum_term;
    grad_term = (y(n) - sigma(n))*X_ext(n, :);
    grad = grad + grad_term;
end
end
```

**Exercise: d)** The script below shows the implementation of gradient ascent (with some learning rate  $\eta$ ) for learning parameters  $\beta$  that maximise the loglikelihood on the training set:


```
function [beta, loglike-track] = train_logreg(X_ext,y, learning_rate)
N = size(X_ext,1); D = size(X_ext,2);
loglike-track = []; i = 1;
%Initialise Beta and grad from Multivariate Gaussian Distribution
grad = mvnrnd(zeros(D,1), eye(D )); beta = mvnrnd(zeros(1, D), eye(D));

while norm(grad) > 0.01 %Expecting grad ~ 0 around a local max
    %Call loglikelihood and its gradient
    [loglike, grad] = log_likelihood_grad(X_train_ext, y_train, beta);
    %Updating beta and number of iterations
    beta = beta + learning_rate * grad;
    %Keeping track of the loglikelihood for later plotting
    loglike-track = [loglike-track, loglike];
    %Convergence control
    if i > 100
        disp('Convergence not achieved after 100 iterations')
        break
    end
    i = i + 1;
end
end
```

The learning rate  $\eta$  has to be small enough, so that the loglikelihood does not oscillate avoiding convergence; and large enough, so that the convergence does not take an infinite amount of time. After some experimentation (shown in the figure below) `learning_rate = 0.005` seems to be the best candidate in terms of number of iterations to convergence.



(a) Gradient ascent with different  $\eta$  values





(b) Loglikelihoods for training and test sets

Alternatively, we can choose a *adaptive* learning rate that adjust its size depending on the grad size and iteration number, but it is beyond the scope of this coursework.

**Exercise: e)** Train the Logistic Classification model on training set `X_train_ext` via:

`train_logreg(X_train_ext,y_train, 0.005)` it returns: `beta = [0.3178; -0.1467; 0.8206]`

The figure below displays the probability contour  $P(y = 1 | \tilde{x}, \hat{\beta})$  (with  $\hat{\beta}$  as our previous MLE of  $\beta$ ) and the original data points. Marked in black the prediction mistakes (false positives and negatives) if a Bayes classifier was used (hard decisions).


 (a) Probability contour using MLE of  $\beta$ 


(b) ROC curves of MLE vs chance classifier

**Exercise: f)** We can compute the final training and test likelihoods per data point obtaining

$$\text{loglike\_train\_pp} = -0.6404, \text{loglike\_test\_pp} = -238.9777$$

Using the script on  $X_{\text{test\_ext}}$  with threshold  $\tau = 0.5$  returns (hard) predictions  $y_{\text{pred}} \in \{0, 1\}$  and the corresponding confusion matrix

```
function [y-pred, confusion_mat] = predicts(y, X_ext, beta, tau)
    N = size(X_ext,1);
    %Predictions based on probabilities P(y = 1 | x, beta)
    probs = probs_logreg(X_ext, beta); %N-column vector with probabilities
    y-pred = probs > tau*ones(N,1); %logic vector
    %Confusion matrix generated by counting frequencies
    true_negatives = sum((y-pred == 0) .* (y == 0));
    false_positives = sum((y-pred == 1) .* (y == 0));
    n_0 = true_negatives + false_positives; % = sum(y_test == 0) = # y-test with y = 0
    false_negatives = sum((y-pred == 0) .* (y == 1));
    true_positives = sum((y-pred == 1) .* (y == 1));
    n_1 = false_negatives + true_positives; % = sum(y_test == 1) = # y-test with y = 1
    confusion_mat = [true_negatives/n_0, false_positives/n_0;
                    false_negatives/n_1, true_positives/n_1];
end
```

Using our previous MLE of  $\beta$  outputs the following confusion matrix:  $\begin{bmatrix} 0.7044 & 0.2956 \\ 0.2932 & 0.7068 \end{bmatrix}$

Alternatively, we can use the Matlab build-in command `confusion()`.

**Exercise: g)** Script below plots the ROC curve of our model (given any extended data set  $X_{\text{ext}}$  and labels  $y$ ) and also computes the area ( $\in [0.5, 1]$ ) under it:

```
function area = roc_curve(X_ext, y, beta)
    x1 = []; x2 = [];
    %Run through possible values of threshold tau
    for tau = 0:0.01:1
        %Hard predictions with threshold tau and corresponding conf_mat
        [y-pred, confusion_mat] = predicts(y, X_ext, beta, tau);
        %Tracking confusion matrix as tau evolves
        x1 = [confusion_mat(1,2), x1]; %false positives
        x2 = [confusion_mat(2,2), x2]; %true positives
    end
end
```

```

end
plot(x1,x2)
hold on
xlabel('False positive rate'); ylabel('True positive rate');
%Area under the ROC Curve using trapezium rule
area = trapz(x1, x2);
end
    
```

For `roc_curve(X_test_ext, y_test, beta)` it returns `area = 0.7314` and the plot (b) in previous page.

**Exercise: h)** We define radial basis functions (RBFs) for given width  $l$ , in the following way:


$$\tilde{x}^{(1)} = 1, \tilde{x}_n^{(m+1)} = \exp\left(\frac{1}{2l^2} \sum_{d=1}^2 (\tilde{x}_n^{(d)} - \tilde{x}_m^{(d)})^2\right) \text{ for } m = 1, \dots, N_{train}$$

Next script computes the extended inputs using RBFs (centred at the training set `X_train` and width  $l$ ) for any data set `X`:


```

function RBFs = RBFs(X, X_train, l)
N = size(X,l); N_train = size(X_train, 1);
%Initialise extend radial basis
RBFs = zeros(N, N_train + 1); %Each input (1,.., N) has N_train + 1 features
%Fill the first column with ones
RBFs(:,1) = ones(N, 1);
%Use loop to fill remaining entries
for m = 1:N_train
    for n = 1:N
        RBFs(n,m+1) = exp(-sum((X(n,:) - X_train(m,:)).^2) / (2*l^2));
    end
end
end
    
```


Now let's train our Logistic Classification model via `train_logreg(RBF,y_train, learning_rate)` using `RBF = RBFs(X_train, X_train, l)` and `y_train` as inputs. The graphs below show the convergence of the loglikelihood flex a fixed value of  $l \in \{0.01, 0.1, 1\}$  and different  $\eta$ -s



(a)  $l = 0.1$



(b)  $l = 0.01$



(c)  $l = 1$

**Exercise: i)** To avoid numerical errors, introduce the following patch in `log_likelihood_grad`:

```

var = X_ext*beta %outside the loop
%...
if sum_term < -5000
    sum_term = (2*y(n)-1)*(var(n));
end
    
```

Where we used the fact that as  $\beta^T \tilde{x}_n \rightarrow -\infty$ ,

$$y_n \log(\sigma(\beta^T \tilde{x}_n)) + (1 - y_n) \log(1 - \sigma(\beta^T \tilde{x}_n)) \sim y_n (\beta^T \tilde{x}_n) + (1 - y_n) (-\beta^T \tilde{x}_n) \sim (2y_n - 1) \beta^T \tilde{x}_n$$

The resulting final training and test loglikelihoods per data point for the different models:

For  $l = 0.01$ , `loglike_train_pp` = -0.1979 `loglike_test_pp` = -265.7482

`area` = 0.5744 with confusion matrix:  $\begin{bmatrix} 0.9759 & 0.0246 \\ 0.8848 & 0.1152 \end{bmatrix}$

For  $l = 0.1$ , `loglike_train_pp` = -0.0482 `loglike_test_pp` = -169.4299

`area` = 0.9163 with confusion matrix:  $\begin{bmatrix} 0.9015 & 0.0985 \\ 0.1466 & 0.8534 \end{bmatrix}$

For  $l = 1$ , `loglike_train_pp` = -0.2235 `loglike_test_pp` = -106.5136

`area` = 0.9529 with confusion matrix:  $\begin{bmatrix} 0.8670 & 0.1330 \\ 0.0838 & 0.9162 \end{bmatrix}$

Both  $l = 1, 0.1$  perform well at predicting the labels for the test set (with `tau` = 0.5), whereas  $l = 0.01$  does not. Remember  $l$  is a measure of the length-scale (or sparseness) of the extended inputs. Hence for a small  $l$  where the inputs are really close to each other we should expect to see soft decision boundaries, giving many wrong predictions. On the other hand, for larger  $l = 1$  we should expect hard decision boundaries

**Exercise: j)** Introduce a Gaussian prior for the parameters  $\beta$ :  $p(\beta^{(m)}) \sim \mathcal{N}(\beta^{(m)}; 0, 1)$  for each of the models in part (i). Then, define the maximum a posteriori (MAP) estimator of  $\beta$  as:

$$\beta_{MAP} := \arg \max_{\beta} p(\beta | X, \mathbf{y}) = \arg \max_{\beta} p(\beta | X) p(\mathbf{y} | X, \beta) \quad (1)$$

$$= \arg \max_{\beta} p(\beta) p(\mathbf{y} | X, \beta) = \arg \max_{\beta} \log(p(\beta)) + \mathcal{L}(\beta) \quad (2)$$

$$= \arg \max_{\beta} \frac{-\beta^T \beta}{2} + \mathcal{L}(\beta) \quad (3)$$

Run gradient ascent for the objective function in line (3) to train our model using the following gradient instead of `log_likelihood_grad`:

```
function grad = posterior_grad(X_ext, y, beta)
N = size(X_ext,1);
sigma = probs_logreg(X_ext, beta);
grad=0;
for n=1:N
    grad_term = (y(n) - sigma(n))*X_ext(n,:);
    grad = grad + grad_term;
end
grad = grad - beta;
end
```


Using the output `beta_map` in `roc_curve` we obtain:

For  $l = 0.01$ , `loglike_train_pp` = -0.1979 `loglike_test_pp` = -265.7482


`area` = 0.5631 with confusion matrix:  $\begin{bmatrix} 0.9901 & 0.0099 \\ 0.8639 & 0.1361 \end{bmatrix}$

For  $l = 0.1$ , `area` = 0.9454 with confusion matrix:  $\begin{bmatrix} 0.9557 & 0.0443 \\ 0.2044 & 0.7966 \end{bmatrix}$


For  $l = 1$ , `area` = 0.9630 with confusion matrix:  $\begin{bmatrix} 0.9113 & 0.0887 \\ 0.0890 & 0.9110 \end{bmatrix}$



(a)  $l = 0.01$



(b)  $l = 0.1$



(c)  $l = 1$

So we see that the MAP for  $\beta$  can also do a great job classifying the test data.